

Charles Haithcock

1 March 2018

CSC 591-604, Hung-Wei Tseng

Homework 1

	Mode 0, Speedup		Mode 1, Speedup	
Run 1	16191837	0	16443246	0.9984
Run 2	16198219	0	16419580	0.9865
Run 3	16200626	0	16417500	0.9867

The threaded mode 1 has slightly worse times than the `qsort()` mode 0. The complexity of the code is inherently $O(n\log(n) + \frac{n}{2}\log(\frac{n}{2}))$ which is thus dominated by the $n\log(n)$. The implementation divides the problem in half and each of the two threads mergesorts the halves. Once complete, a final merge is done on the entire dataset, as, upon completion, the data consists of two sorted subarrays. The final merge provides the cost of $n\log(n)$ while each of the threads takes on half the input and thus $\frac{n}{2}\log(\frac{n}{2})$. As the work is divided into two processes for each processor to work on, the complexity of the threads is not $n\log(\frac{n}{2})$. The complexity is difficult to compare with the C `qsort` function as there is no clear description within the standard C library ISO or otherwise defining the complexity.

Considering my implementation of parallelized mergesort performs almost as well as `qsort` provided from the standard C library, improvements should be made. For example, given time, the temporary storage buffer for subarrays could simply be switched off with the real data input while sorting rather than `memcpy()` every time a subarray was sorted. Additional threads were attempted (up to 6 were attempted) but the performance was best at 2 threads. I had no board mates so no way to compare the run times.